

Exercise 5: Practice with MLP Network (3)

Assignment 4 Author: David Nogales Pérez

```
In [ ]: import numpy as np
import pandas as pd
import seaborn as sns
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.activations import sigmoid
from tensorflow.keras.losses import MeanSquaredError

from sklearn.model_selection import RepeatedKFold, cross_val_score
from tensorflow.keras.wrappers.scikit_learn import KerasRegressor
from tensorflow.keras.metrics import RootMeanSquaredError
from keras import losses
from keras import regularizers
import keras
import numpy as np
import matplotlib.pyplot as plt
import warnings
warnings.filterwarnings('ignore')
```

Generate training data

```
In [ ]: def y1(k,y,u):
    pi = np.pi
    e = np.e
    u_k = u[k]
    y_k = y[k]
    exp = -(u_k**2)-(y_k**2)
    return 2.5*y_k*np.sin(Pi*np.power(e,exp))
def y2(k,y,u):
    u_k=u[k]
    return u_k*(1+u_k**2)

#Generates n samples for the function y(k) described in the problem statement
def generate_data(n=10,seed=12345):
    data = np.zeros(n)
    rng = np.random.default_rng(seed)
    u = rng.uniform(-2,2,n)
    for i in range(1,n):
        yk=y1(i-1,data,u)+y2(i-1,data,u)
        data[i]=yk
    return np.arange(n),data

n = 500 #training size
X_train,y_train = generate_data(n)
X_train.shape,y_train.shape

Out[ ]: ((500,), (500,))

In [ ]: plt.scatter(X_train,y_train)
plt.title(f'Function y(k) for n={n}')
plt.ylabel("y(k)")
plt.xlabel("k")
plt.show()
```

Plot of 500 samples of the function:

where

$$y(k) = y_1(k-1) + y_2(k-1)$$
$$y_1(k) = 2.5y(k)\sin(\pi e^{-u^2(k)-y^2(k)})$$
$$y_2(k) = u(k)(1 + u^2(k))$$

Build MLP

```
In [ ]: #Generates the Loss plot for the given history generated by the model
def plot_loss(history):
    plt.plot(history.history['loss'])
    plt.title('model loss')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()

#Generates a scatter plot of the input data and the prediction of the given model
def plot_predicted(model,X,y):
    plt.scatter(X,y,label="real")
    plt.scatter(X,model.predict(X),label="predicted")
    plt.legend(loc="upper right")
    plt.xlabel("k")
    plt.ylabel("Value")

#Returns a function which generates a 3-hidden Layer MLP with the given regularization and activation function
def buildmodel(act='sigmoid',reg = regularizers.l1_l2()):
    def build():
        model = Sequential()
        model.add(Dense(1, input_dim=1, activation=act, kernel_regularizer=reg))
        model.add(Dense(50, activation=act, kernel_regularizer=reg))
        model.add(Dense(30, activation=act, kernel_regularizer=reg))
        model.add(Dense(15, activation=act, kernel_regularizer=reg))
        model.add(Dense(1,activation = 'linear'))
        model.compile(optimizer="Adam", loss=losses.mean_squared_error,metrics=[RootMeanSquaredError()])
        return model
    return build
```

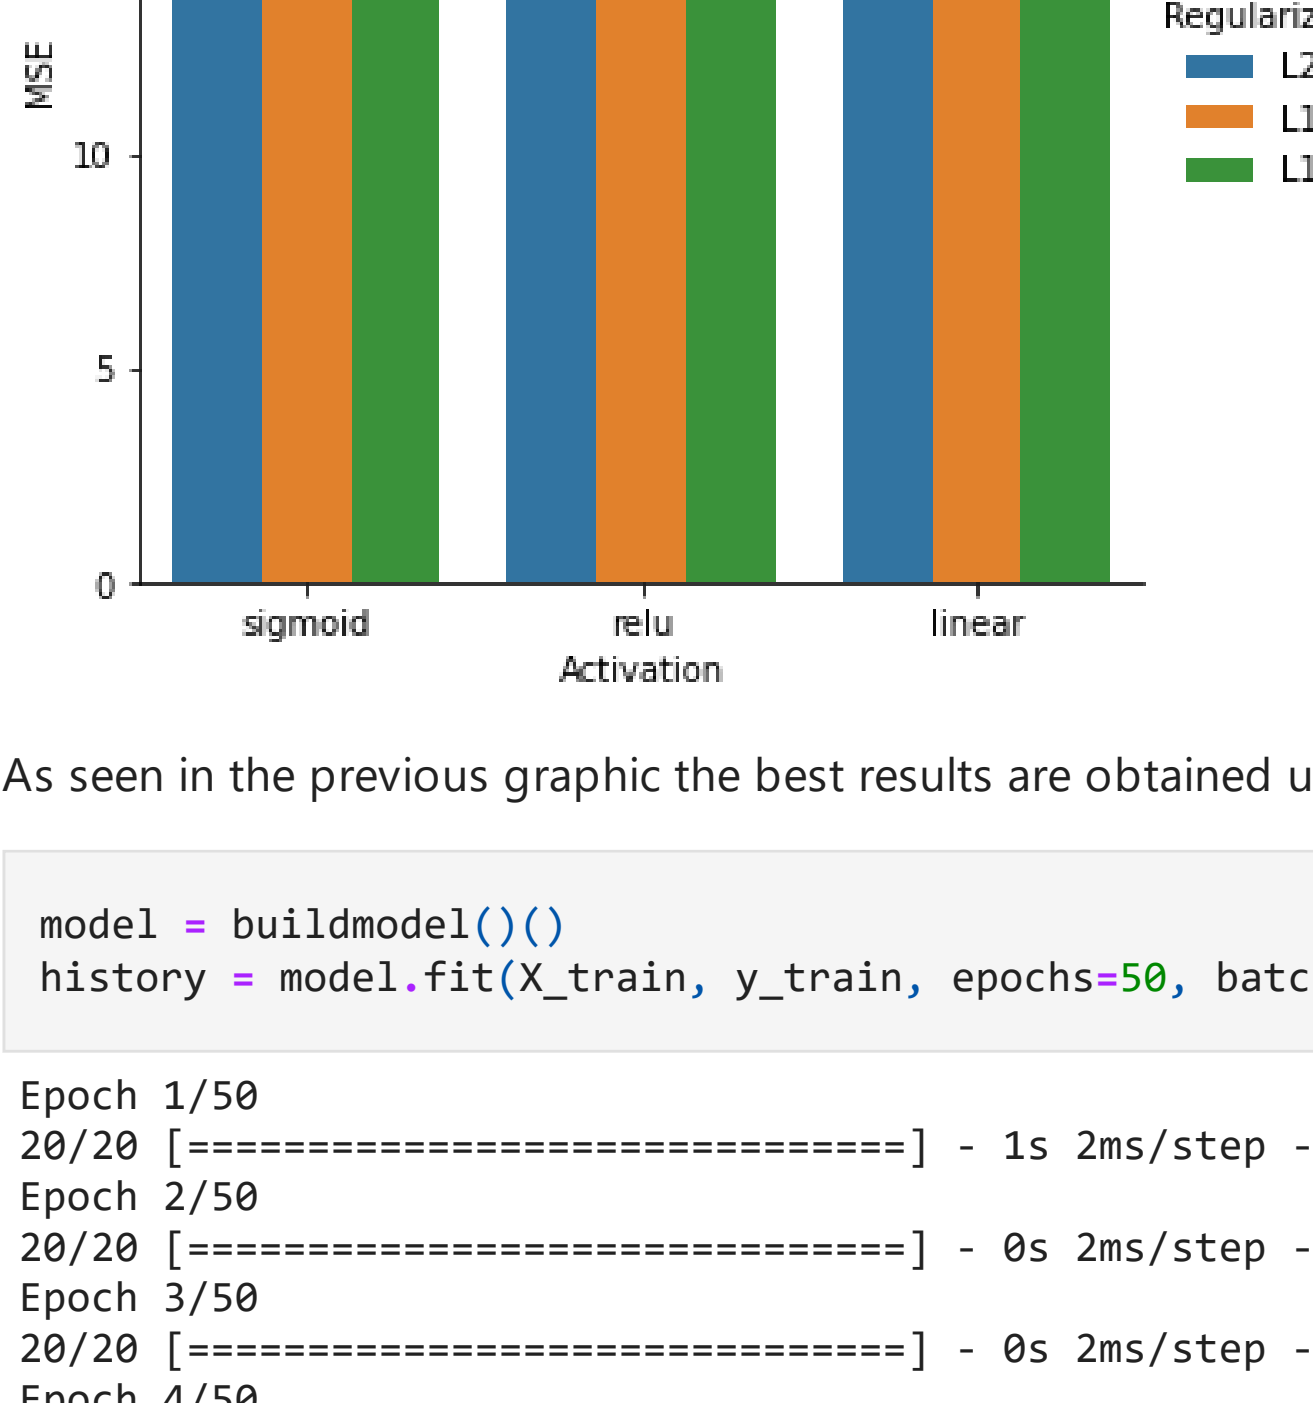
Applying cross validation to the MLP

```
In [ ]: parameters = {"activation":["sigmoid","linear","relu"],
                    "regularizer":[(("L1",regularizers.l1()),("L2",regularizers.l2()),("L1_L2",regularizers.l1_l2()))]
}

results = {
}
for act in parameters["activation"]:
    results[act]={}
    for reg in parameters["regularizer"]:
        build = buildmodel(act,reg[1])
        estimator= KerasRegressor(build_fn=build, epochs=50, batch_size=25, verbose=0)
        kfolds= RepeatedKFold(n_splits=5, n_repeats=5)
        result = cross_val_score(estimator, X_train, y_train, cv=kfolds, n_jobs=1)
        mean = result.mean()
        print(f"Cross validation mean of MSE: {mean}")
        results[act][reg[0]] = mean

Cross validation mean of MSE: -16.952189750671387
Cross validation mean of MSE: -16.827399063110352
Cross validation mean of MSE: -16.93697811126709
Cross validation mean of MSE: -22.822959213256837
Cross validation mean of MSE: -18.354305610656738
Cross validation mean of MSE: -22.08200454711914
Cross validation mean of MSE: -17.559741706848143
Cross validation mean of MSE: -16.844078407287597
Cross validation mean of MSE: -17.384937133789062
```

```
In [ ]: df = pd.DataFrame(results)
df = df.rename_axis('Regularization').reset_index()
df = pd.melt(df,id_vars=["Regularization"],value_vars=['sigmoid',"linear","relu"],var_name='Activation',value_name='MSE').sort_values('MSE',ascending=False).reset_index(drop=True)
df['MSE']=abs(df['MSE'])
ax = sns.catplot(data=df,x='Activation',y='MSE',hue='Regularization',kind='bar').set(title = 'MSE Metrics For Each Activation Function and Regularization')
df
```

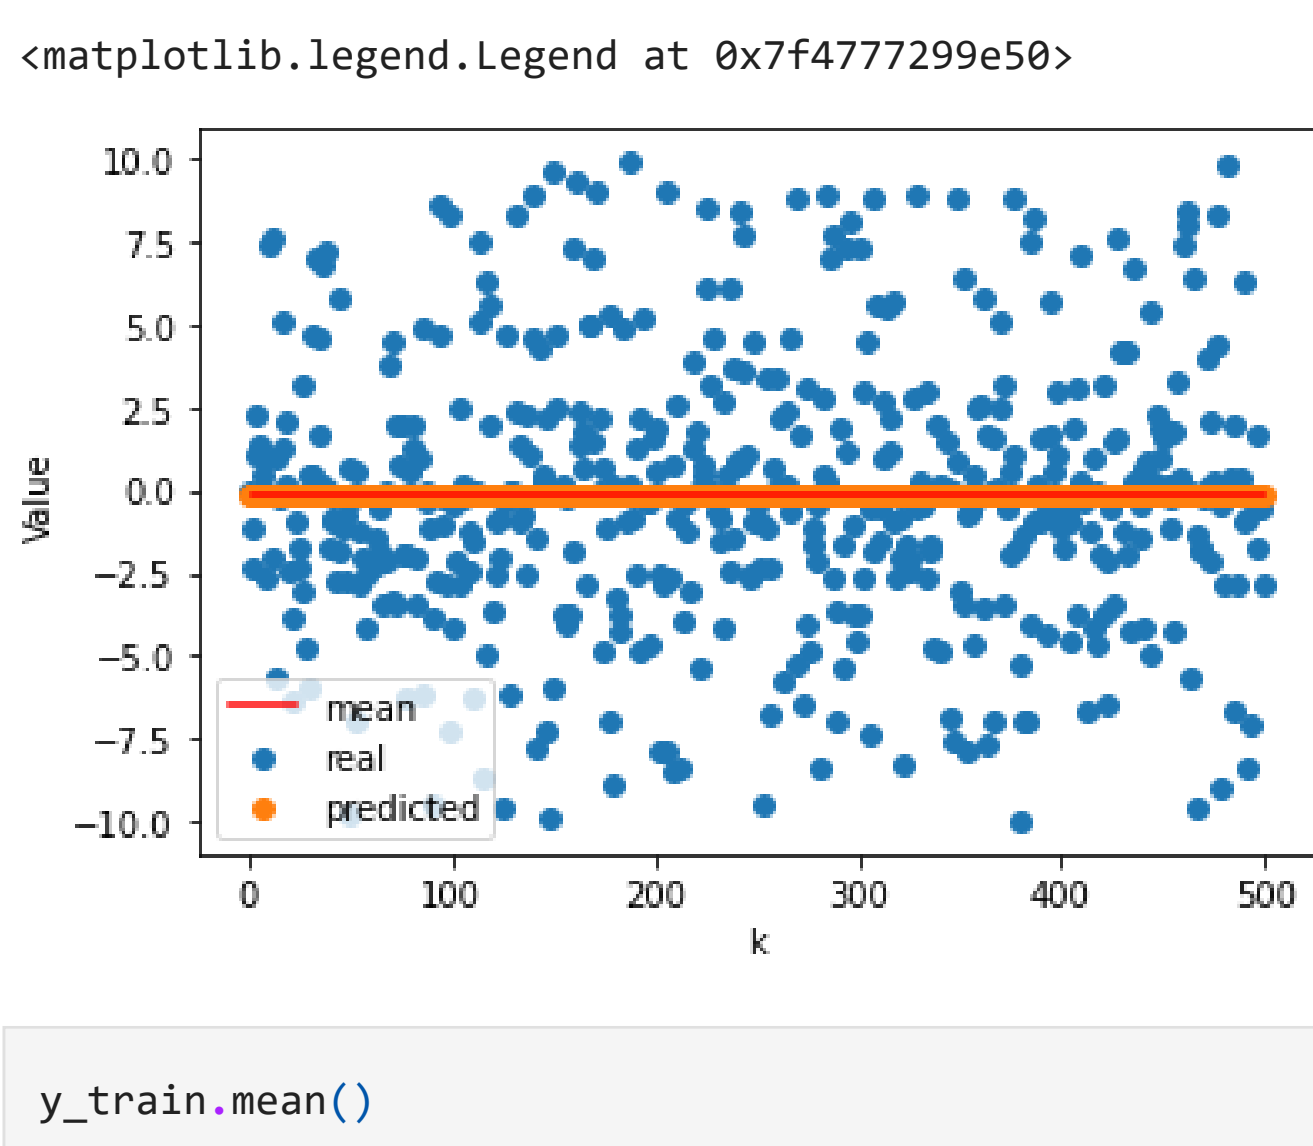


As seen in the previous graphic the best results are obtained using a Sigmoid activation function with any kind of regularization but L1+L2 regularization will be used from now on.

```
In [ ]: model = buildmodel()()
history = model.fit(X_train, y_train, epochs=50, batch_size=25, verbose=1)

Epoch 1/50 [=====] - 1s 2ms/step - loss: 20.2313 - root_mean_squared_error: 4.0947
Epoch 2/50 [=====] - 0s 2ms/step - loss: 19.7954 - root_mean_squared_error: 4.0811
Epoch 3/50 [=====] - 0s 2ms/step - loss: 19.4505 - root_mean_squared_error: 4.0776
Epoch 4/50 [=====] - 0s 2ms/step - loss: 19.1481 - root_mean_squared_error: 4.0771
Epoch 5/50 [=====] - 0s 2ms/step - loss: 18.8636 - root_mean_squared_error: 4.0764
Epoch 6/50 [=====] - 0s 2ms/step - loss: 18.6112 - root_mean_squared_error: 4.0767
Epoch 7/50 [=====] - 0s 2ms/step - loss: 18.3772 - root_mean_squared_error: 4.0761
Epoch 8/50 [=====] - 0s 2ms/step - loss: 18.1737 - root_mean_squared_error: 4.0767
Epoch 9/50 [=====] - 0s 2ms/step - loss: 17.9868 - root_mean_squared_error: 4.0764
Epoch 10/50 [=====] - 0s 2ms/step - loss: 17.8301 - root_mean_squared_error: 4.0770
Epoch 11/50 [=====] - 0s 2ms/step - loss: 17.6890 - root_mean_squared_error: 4.0775
Epoch 12/50 [=====] - 0s 2ms/step - loss: 17.5566 - root_mean_squared_error: 4.0767
Epoch 13/50 [=====] - 0s 2ms/step - loss: 17.4756 - root_mean_squared_error: 4.0788
Epoch 14/50 [=====] - 0s 2ms/step - loss: 17.3640 - root_mean_squared_error: 4.0760
Epoch 15/50 [=====] - 0s 2ms/step - loss: 17.2929 - root_mean_squared_error: 4.0763
Epoch 16/50 [=====] - 0s 2ms/step - loss: 17.2348 - root_mean_squared_error: 4.0762
Epoch 17/50 [=====] - 0s 2ms/step - loss: 17.1856 - root_mean_squared_error: 4.0760
Epoch 18/50 [=====] - 0s 2ms/step - loss: 17.1455 - root_mean_squared_error: 4.0764
Epoch 19/50 [=====] - 0s 2ms/step - loss: 17.1065 - root_mean_squared_error: 4.0760
Epoch 20/50 [=====] - 0s 2ms/step - loss: 17.0780 - root_mean_squared_error: 4.0764
Epoch 21/50 [=====] - 0s 2ms/step - loss: 17.0525 - root_mean_squared_error: 4.0764
Epoch 22/50 [=====] - 0s 2ms/step - loss: 17.0300 - root_mean_squared_error: 4.0766
Epoch 23/50 [=====] - 0s 2ms/step - loss: 17.0010 - root_mean_squared_error: 4.0757
Epoch 24/50 [=====] - 0s 2ms/step - loss: 16.9856 - root_mean_squared_error: 4.0764
Epoch 25/50 [=====] - 0s 2ms/step - loss: 16.9640 - root_mean_squared_error: 4.0758
Epoch 26/50 [=====] - 0s 2ms/step - loss: 16.9537 - root_mean_squared_error: 4.0763
Epoch 27/50 [=====] - 0s 2ms/step - loss: 16.9360 - root_mean_squared_error: 4.0759
Epoch 28/50 [=====] - 0s 2ms/step - loss: 16.9286 - root_mean_squared_error: 4.0766
Epoch 29/50 [=====] - 0s 2ms/step - loss: 16.9162 - root_mean_squared_error: 4.0764
Epoch 30/50 [=====] - 0s 2ms/step - loss: 16.8990 - root_mean_squared_error: 4.0758
Epoch 31/50 [=====] - 0s 2ms/step - loss: 16.8949 - root_mean_squared_error: 4.0764
Epoch 32/50 [=====] - 0s 2ms/step - loss: 16.8891 - root_mean_squared_error: 4.0769
Epoch 33/50 [=====] - 0s 2ms/step - loss: 16.8755 - root_mean_squared_error: 4.0765
Epoch 34/50 [=====] - 0s 2ms/step - loss: 16.8672 - root_mean_squared_error: 4.0766
Epoch 35/50 [=====] - 0s 2ms/step - loss: 16.8631 - root_mean_squared_error: 4.0771
Epoch 36/50 [=====] - 0s 2ms/step - loss: 16.8439 - root_mean_squared_error: 4.0759
Epoch 37/50 [=====] - 0s 2ms/step - loss: 16.8511 - root_mean_squared_error: 4.0773
Epoch 38/50 [=====] - 0s 2ms/step - loss: 16.8461 - root_mean_squared_error: 4.0775
Epoch 39/50 [=====] - 0s 2ms/step - loss: 16.8436 - root_mean_squared_error: 4.0778
Epoch 40/50 [=====] - 0s 2ms/step - loss: 16.8231 - root_mean_squared_error: 4.0766
Epoch 41/50 [=====] - 0s 2ms/step - loss: 16.8064 - root_mean_squared_error: 4.0757
Epoch 42/50 [=====] - 0s 2ms/step - loss: 16.8026 - root_mean_squared_error: 4.0761
Epoch 43/50 [=====] - 0s 2ms/step - loss: 16.8018 - root_mean_squared_error: 4.0766
Epoch 44/50 [=====] - 0s 2ms/step - loss: 16.7991 - root_mean_squared_error: 4.0767
Epoch 45/50 [=====] - 0s 2ms/step - loss: 16.7844 - root_mean_squared_error: 4.0758
Epoch 46/50 [=====] - 0s 2ms/step - loss: 16.7790 - root_mean_squared_error: 4.0759
Epoch 47/50 [=====] - 0s 2ms/step - loss: 16.7776 - root_mean_squared_error: 4.0763
Epoch 48/50 [=====] - 0s 2ms/step - loss: 16.7693 - root_mean_squared_error: 4.0759
Epoch 49/50 [=====] - 0s 2ms/step - loss: 16.7808 - root_mean_squared_error: 4.0776
Epoch 50/50 [=====] - 0s 2ms/step - loss: 16.7632 - root_mean_squared_error: 4.0762

In [ ]: plot_loss(history)
```



Observing the loss curve we can see that the NN is starting to overfit slightly so no extra epoch are required to fully train the NN.

```
In [ ]: plot_predicted(model,X_train,y_train)
plt.plot(np.arange(0,n)+1,np.zeros(n)y_train.mean(),'-r',label="mean")
plt.legend()
```

Out[]: <matplotlib.legend.Legend at 0x7f4777299e50>

In []: y_train.mean()

Out[]: -0.10853680066900842

In this plot we can see how the MLP predicts a steady line avoiding the noise introduced by the randomness of $u(k)$ in the system. Which is really close to the mean of the image values of the function.